PROGRAMMING ASSIGNMENT - 2

Under the guidance of Prof. Ioan Raicu

Abhay Nagaraj B R

Indranil Thakur

PROGRAMMING ASSIGNMENT - 2

Problem statement:

Perform sorting of large and moderate datasets in different configurations and compare their performances.

Methodology:

Implement the sorting of data using Merge Sort which seemed much feasible for the current scenario.
Vary the AWS ec2 instance type and the number of instances used to form the cluster to record the speedup correspondingly.
Gensort is used to generate the data on the fly instead of wasting time and money on the transfer of dat to be sorted.
Valsort is used to validate the output of the sort algorithm used.

Runtime Environment settings needed are as follows:

- Virtual cluster setup with automation scripts: 10 points
- • Required: Shared memory terasort implementation/scripts: 30 points
- • Hadoop Terasort implementation/scripts: 10 points
- • Spark Terasort implementation/scripts: 10 points
- • MPI Terasort implementation/scripts: 10 points
- • Readme.txt: 5 points
- • Performance evaluation, data, explanations, etc: 15 points per system (for a maximum of 60 points)
- • Followed instructions on deliverables: 5 points

This programming assignment has following parts

- Shared Memory Sort
- Hadoop Pseudo Mode Sorting
- Hadoop 8 nodes cluster Sorting
- Spark Pseudo Mode Sorting
- Spark 8 node cluster Sorting

The assignment had 8 parts as explained below

Virtual cluster setup with automation scripts:

Pre requisites:
An AWS account to be created.
For easier handling of access keys and .pem files on windows PCs (which is what I have used) -
Install putty/puttyGen
Install WinSCP for copying and editing of files on the virtual nodes with a free flowing UI.
This section is common for all the other configurations where the AWS cluster has to be set up either **pseudo** or **multi-node.**

Steps:

- Launch an ec2 instance.

- Create and download the keypair.pem file on your local machine.
- Use puttyGen to generate a private key of .ppk type.
- Use the above generated .ppk file in WinSCP to connect to the instances

Now we need to set up a password-less authentication for which we have created a config file which looks like the following.

It contains the host alias name and its public DNS address.
The username to be used.
The location of the .pem file on the virtual node.

```
Host namenode
  HostName ec2-54-152-0-85.compute-1.amazonaws.com
  User ubuntu
  IdentityFile ~/.ssh/my-hadoop-key.pem
Host datanode1
  HostName ec2-34-201-116-73.compute-1.amazonaws.com
  User ubuntu
  IdentityFile ~/.ssh/my-hadoop-key.pem
Host datanode2
  HostName ec2-54-211-103-104.compute-1.amazonaws.com
  User ubuntu
  IdentityFile ~/.ssh/my-hadoop-key.pem
Host datanode3
  HostName ec2-52-23-195-185.compute-1.amazonaws.com
  User ubuntu
  IdentityFile ~/.ssh/my-hadoop-key.pem
```

These .pem files are copied to all the nodes in the cluster, so that the nodes can talk to each other through a password-less mode of authentication.

```
scp ~/.ssh/my-hadoop-key.pem ~/.ssh/config datanode1:~/.ssh
```

Then we use the ssh-keygen to generate a public key on the virtual nodes.

```
#generate key file on the one of the nodes
 ssh-keygen -f ~/.ssh/ssh_rsa -t rsa -P ""
 -f : file name which follows it
 -t : type of key which is rsa here
 -P : passphrase which is "" empty here
```

The content of this .pub file generated is copied into the authorized_keys file in all the nodes in the cluster.

Now that we get the access the cluster nodes, Java, Hadoop and Spark are installed.

Java version installed – 1.8

Hadoop version installed – 2.7.4

Setting the environment variables
.profile file in every node is modified as follows:

```
#set environment variables in all the nodes
export JAVA_HOME=/usr
export PATH=$PATH:$JAVA_HOME/bin
export HADOOP_HOME=/usr/local/hadoop
export PATH=$PATH:$HADOOP_HOME/bin
export HADOOP_CONF_DIR=/usr/local/hadoop/etc/hadoop
```

The modified .profile is loaded to bring it into effect

```
#load variables
. ~/.profile
```

Hadoop Configuration is also changed on all the nodes.

```
#hadoop configuration on all nodes
#$HADOOP_CONF_DIR/hadoop-env.sh change JAVA_HOMEexport
JAVA_HOME=/usr
```

Virtual Cluster (1-node i3.large):

Shared memory

1-node 13.large instance

Java has been installed on the instance. On demand instance was used.

Instance Specifications

Instance ID i-0e6add78e8b5eb12f	Public DNS (IPv4) ec2-35-153-168-54.compute-1.amazonaws.com
Instance state Running	IPv4 Public IP 35.153.168.54
Instance type i3.large	
Private DNS ip-172-31-33-19.ec2.internal	
Availability zone us-east-1c	Private IPs 172.31.33.19

Case 1: Data to be sorted: 128 GB

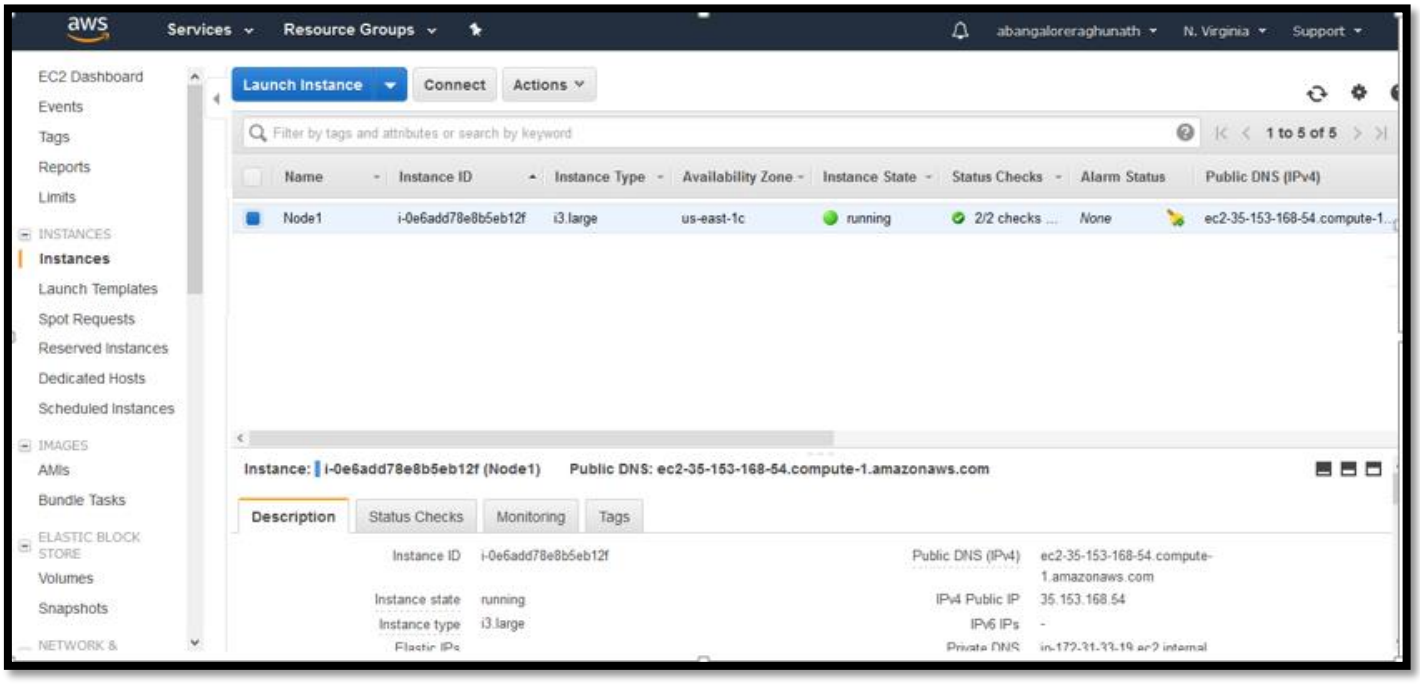

Since the instance i3.xlarge itself has a memory of 16GB and an SSD of 440GB attached to it, we don't need a RAID0 to be configured for this instance.

Therefore the code is deployed on the node and is run to get following output. We have used 4 threads.

It completed in 3 hours and 40 minutes

Hadoop

128GB

Hadoop 2.7.4 was the installed version.
Difficulty faced: When we tried installing Java 1.7 it did not work with this version of Hadoop and were stuck in this stage for a long time. Later on after a long search, we understood that this version of hadoop is not compatible with Java 1.7. It needs java 1.8 or newer.

Furthermore, Hadoop MapReduce job needs a few configurations to be made before hand running the job.

The following is the config file.

```
#NameNode configuration
#HADOOP_CONF_DIR/Core-site.xml change configuration element
#change the namenode_public_dns to your NameNode Public DNS
<configuration>
        <property>
                <name>fs.defaultFS</name>
                <value>hdfs://namenode_public_dns:9000</value>
        </property>
</configuration>

#$HADOOP_CONF_DIR/yarn-site.xml change configuration element
#change the namenode_public_dns to your NameNode Public DNS
<configuration>
        <property>
                <name>dfs.replication</name>
                <value>1</value>
        </property>
        <property>
                <name>dfs.namenode.name.dir</name>
                <value>file:///usr/local/hadoop/etc/hadoop_data/hdfs/namenode</value>
        </property>
</configuration>
```

In the above configuration setting, especially in the yarn-site.xml, we have set the **replication factor to 1** Which means that the data blocks are copied only once to the datanodes. Therefore, we lose on the resilience of the cluster from any failures, but we do reduce the storage requirement and save time from writing multiple copies of data to blocks.

Below are the environment variable settings needed to use aliases on the bash command line.

#set environment variables in all the nodes
export JAVA_HOME=/usr
export PATH=$PATH:$JAVA_HOME/bin
export HADOOP_HOME=/usr/local/hadoop
export PATH=$PATH:$HADOOP_HOME/bin
export HADOOP_CONF_DIR=/usr/local/hadoop/etc/hadoop

```
[edureka@localhost jobJars]$ yarn jar CounterDemo.jar com.edureka.mapreduce.job.WordCountJob /user/edureka/wordcountproblemm /user/edureka/WC1
17/12/04 07:38:05 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
17/12/04 07:38:05 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
17/12/04 07:38:06 INFO input.FileInputFormat: Total input paths to process : 1
17/12/04 07:38:06 INFO mapreduce.JobSubmitter: number of splits:2
17/12/04 07:38:06 INFO Configuration.deprecation: user.name is deprecated. Instead, use mapreduce.job.user.name
17/12/04 07:38:06 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar
17/12/04 07:38:06 INFO Configuration.deprecation: mapred.output.value.class is deprecated. Instead, use mapreduce.job.output.value.class
17/12/04 07:38:06 INFO Configuration.deprecation: mapreduce.map.class is deprecated. Instead, use mapreduce.job.map.class
17/12/04 07:38:06 INFO Configuration.deprecation: mapred.job.name is deprecated. Instead, use mapreduce.job.name
17/12/04 07:38:06 INFO Configuration.deprecation: mapreduce.reduce.class is deprecated. Instead, use mapreduce.job.reduce.class
17/12/04 07:38:06 INFO Configuration.deprecation: mapreduce.inputformat.class is deprecated. Instead, use mapreduce.job.inputformat.class
17/12/04 07:38:06 INFO Configuration.deprecation: mapred.input.dir is deprecated. Instead, use mapreduce.input.fileinputformat.inputdir
17/12/04 07:38:06 INFO Configuration.deprecation: mapred.output.dir is deprecated. Instead, use mapreduce.output.fileoutputformat.outputdir
17/12/04 07:38:06 INFO Configuration.deprecation: mapreduce.outputformat.class is deprecated. Instead, use mapreduce.job.outputformat.class
17/12/04 07:38:06 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
17/12/04 07:38:06 INFO Configuration.deprecation: mapred.output.key.class is deprecated. Instead, use mapreduce.job.output.key.class
17/12/04 07:38:06 INFO Configuration.deprecation: mapred.working.dir is deprecated. Instead, use mapreduce.job.working.dir
17/12/04 07:38:07 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1496733246656_0058
17/12/04 07:38:07 INFO impl.YarnClientImpl: Submitted application application_1496733246656_0058 to ResourceManager at /0.0.0.0:8032
17/12/04 07:38:07 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1496733246656_0058/
17/12/04 07:38:07 INFO mapreduce.Job: Running job: job_1496733246656_0058
17/12/04 07:38:14 INFO mapreduce.Job: Job job_1496733246656_0058 running in uber mode : false
17/12/04 07:38:15 INFO mapreduce.Job:  map 0% reduce 0%
17/12/04 07:38:29 INFO mapreduce.Job:  map 17% reduce 0%
17/12/04 07:38:32 INFO mapreduce.Job:  map 24% reduce 0%
17/12/04 07:38:35 INFO mapreduce.Job:  map 25% reduce 0%
17/12/04 07:38:56 INFO mapreduce.Job:  map 40% reduce 0%
17/12/04 07:39:00 INFO mapreduce.Job:  map 44% reduce 0%
17/12/04 07:39:24 INFO mapreduce.Job:  map 50% reduce 0%
17/12/04 07:39:27 INFO mapreduce.Job:  map 53% reduce 0%
17/12/04 07:39:30 INFO mapreduce.Job:  map 56% reduce 0%
```

```
                Map output records=28006216
                Map output bytes=290097360
                Map output materialized bytes=346109804
                Input split bytes=234
                Combine input records=0
                Combine output records=0
                Reduce input groups=355
                Reduce shuffle bytes=346109804
                Reduce input records=28006216
                Reduce output records=355
                Spilled Records=84018648
                Shuffled Maps =2
                Failed Shuffles=0
                Merged Map outputs=2
                GC time elapsed (ms)=4174
                CPU time spent (ms)=17688000
                Physical memory (bytes) snapshot=454291456
                Virtual memory (bytes) snapshot=1082220544
                Total committed heap usage (bytes)=352985088
        Shuffle Errors
                BAD_ID=0
                CONNECTION=0
                IO_ERROR=0
                WRONG_LENGTH=0
                WRONG_MAP=0
                WRONG_REDUCE=0
        File Input Format Counters
                Bytes Read=135061500164
        File Output Format Counters
                Bytes Written=52758399
SUCCESS
```

It completed in 4.5 hours = 17,688 seconds

Spark

128GB

```
ubuntu@ip-172-31-7-126:~$ cd /opt
ubuntu@ip-172-31-7-126:/opt$ ls
spark  spark-2.2.0-bin-hadoop2.7
ubuntu@ip-172-31-7-126:/opt$ spark-submit --version
Welcome to
      ____              __
     / __/__  ___ _____/ /__
    _\ \/ _ \/ _ `/ __/  '_/
   /___/ .__/\_,_/_/ /_/\_\   version 2.2.0
      /_/

Using Scala version 2.11.8, OpenJDK 64-Bit Server VM, 1.8.0_151
Branch
Compiled by user jenkins on 2017-06-30T22:58:04Z
Revision
Url
Type --help for more information.
ubuntu@ip-172-31-7-126:/opt$
```

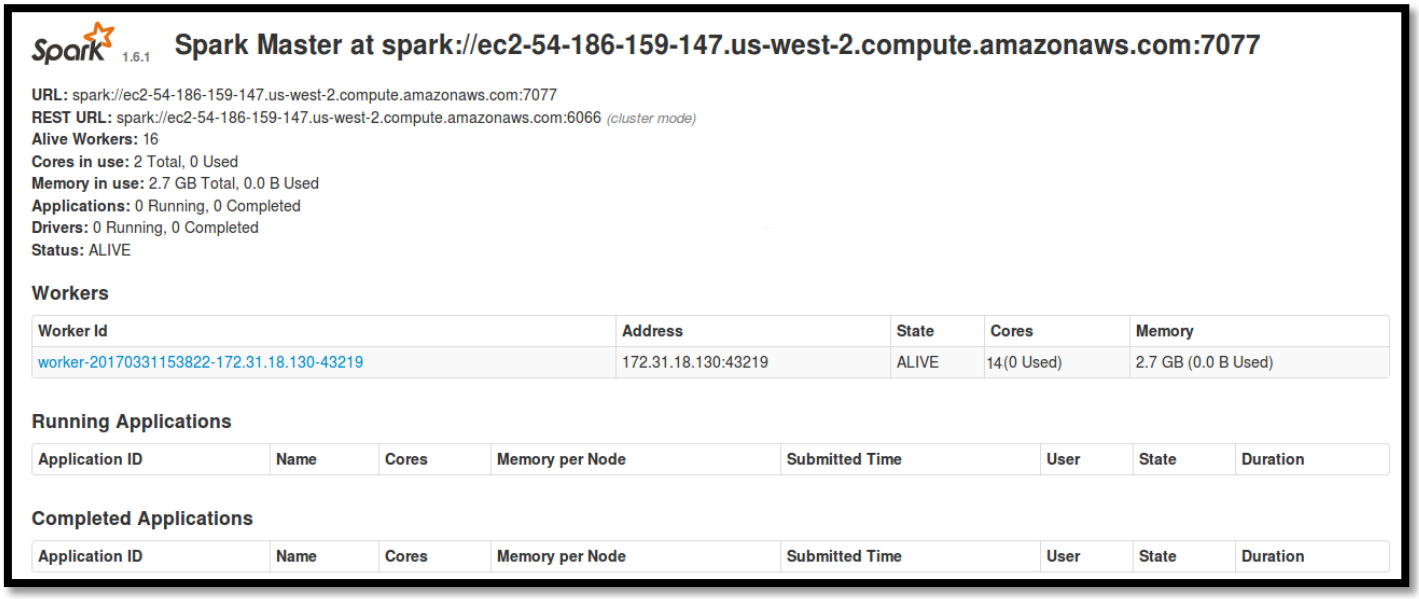

Spark 1.6.1 **Spark Master at spark://ec2-54-186-159-147.us-west-2.compute.amazonaws.com:7077**

URL: spark://ec2-54-186-159-147.us-west-2.compute.amazonaws.com:7077
REST URL: spark://ec2-54-186-159-147.us-west-2.compute.amazonaws.com:6066 *(cluster mode)*
Alive Workers: 16
Cores in use: 2 Total, 0 Used
Memory in use: 2.7 GB Total, 0.0 B Used
Applications: 0 Running, 0 Completed
Drivers: 0 Running, 0 Completed
Status: ALIVE

Workers

Worker Id	Address	State	Cores	Memory
worker-20170331153822-172.31.18.130-43219	172.31.18.130:43219	ALIVE	14(0 Used)	2.7 GB (0.0 B Used)

Running Applications

Application ID	Name	Cores	Memory per Node	Submitted Time	User	State	Duration

Completed Applications

Application ID	Name	Cores	Memory per Node	Submitted Time	User	State	Duration

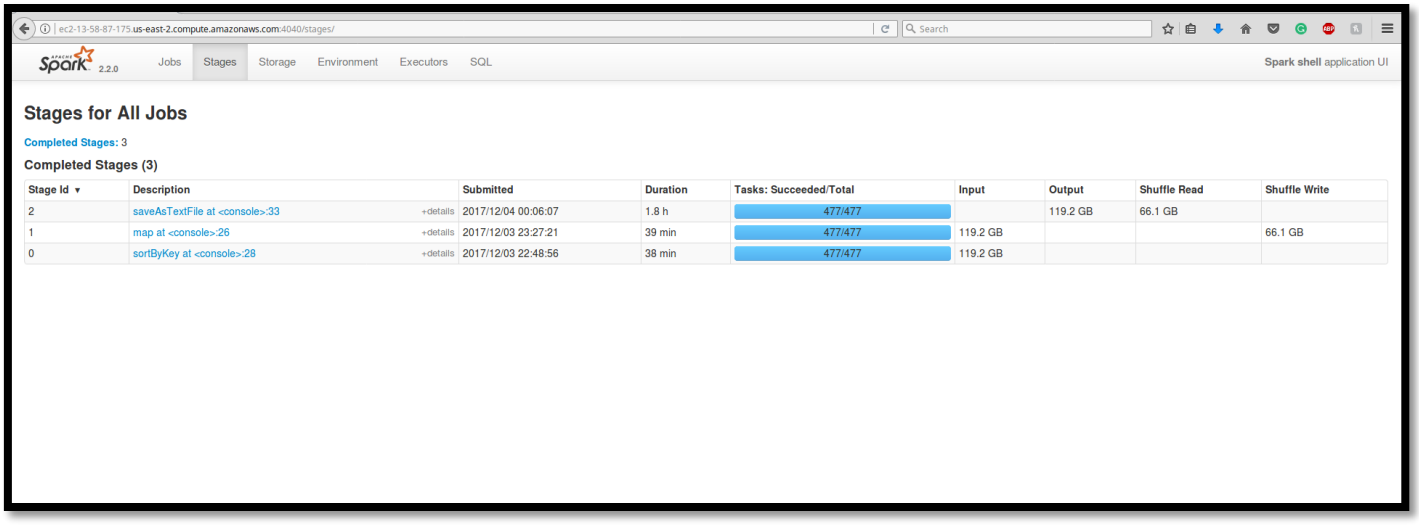

ec2-13-58-87-175.us-east-2.compute.amazonaws.com:4040/stages/

Spark 2.2.0 Jobs | Stages | Storage | Environment | Executors | SQL — Spark shell application UI

Stages for All Jobs

Completed Stages: 3

Completed Stages (3)

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
2	saveAsTextFile at <console>:33 +details	2017/12/04 00:06:07	1.8 h	477/477		119.2 GB	66.1 GB	
1	map at <console>:26 +details	2017/12/03 23:27:21	39 min	477/477	119.2 GB			66.1 GB
0	sortByKey at <console>:28 +details	2017/12/03 22:48:56	38 min	477/477	119.2 GB			

It completed in 184 minutes which is way faster than the Hadoop and the shared memory.

```
root@ip-172-31-24-63 output]$ ./valsort part-00000
-bash: ./valsort: No such file or directory
root@ip-172-31-24-63 output]$ /root/myfile/64/./valsort part-00000
Records: 1316708
Checksum: a0bac38f0ff04
Duplicate keys: 0
SUCCESS - all records are in order
root@ip-172-31-24-63 output]$ /root/myfile/64/./valsort part-00079
Records: 1084439
Checksum: 845a4050be0e7
Duplicate keys: 0
SUCCESS - all records are in order
root@ip-172-31-24-63 output]$ head part-00000
    "O!uve  00000000000000000000000001228D4  77778888000022224444DDDDDDDDEEEE00000000CCCC7777DDDD
    PMd32=  00000000000000000000000003440CC1  FFFFEEEE6666CCCCBBBB9999333335555DDDDDDDD777788886666
    ^3CO],  00000000000000000000000158C5C5  5555AAAA9999EEEE888822229999CCCCDDDD6666555544442222
   !&S3/]]  00000000000000000000000002145D78  8888BBBBDDDD1111CCCC55556666BBBB1111EEEEDDDD22229999
   !,=U#,9  00000000000000000000000019072E3  33332222FFFFBBBB0000FFFFAAAA6666555533333DDDD3333CCCC
   !0f[ITd  00000000000000000000000003CAAB4B  9999FFFF555533337777CCCC4444BBBB7777EEEEBBBBDDDD4444
   !f6Suy2  00000000000000000000000003ABFD84  EEEE555555556666AAAA5555BBBBDDDD000011116666000DDDD
   #%NIpq.  00000000000000000000000003B36FB9  11110000333344441111666666666AAAAAAAA00001111CCCCEEEE
   #'^cl'~  00000000000000000000000002EDC5C8  8888AAAA11114444FFFF77773333EEEE44440000FFFF99999999
   $"-'Q)]  00000000000000000000000005F1265D  CCCC6666EEEE22220000DDDDAAAA88886666BBBB00006666AAAA
root@ip-172-31-24-63 output]$ tail part-00079
~~~uq2k#=U  00000000000000000000000002C06745  99991111DDDD222211110000FFFFEEEEFFFF33337777CCCC2222
~~~v/0&Qnm  00000000000000000000000004709701  CCCC88883333FFFF000000000000999911111FFFF777744446666
~~~yKOl:gE  00000000000000000000000002048B4F  CCCC11114444888822226666BBBB888855557777EEEEBBBB0000
~~~yK^H.il  00000000000000000000000000463D004  44440000FFFF33339999444477777DDDDFFFFAAAA11118888DDDD
~~~yL;C'XE  00000000000000000000000005B0D211  2222EEEE3333000022221111CCCCFFFF555577774444BBBB6666
~~~zbA_ Tt  00000000000000000000000007F9F4F  BBBBCCCC6666555599999FFFF8888AAAA11116666AAAABBBB0000
~~~zeO^FEg  00000000000000000000000001E06130  4444CCCCBBBB999922228888555588888CCCCFFFF000011111111
~~~}GxjWHI  00000000000000000000000000CA1345  777711118888AAAAAAAA22221111BBBB00002222BBBBCCCC2222
~~~}P;]g0g  00000000000000000000000000040DA3E4  4444FFFF4444666663333EEEE88888888DDDDEEEE44442222DDDD
~~~}kU|K<p  00000000000000000000000005E4A0AA  0000666655551111BBBB88889999AAAA55550000333355557777
root@ip-172-31-24-63 output]$
```

The output was verified using Valsort as above.

Virtual Cluster (1-node i3.4xlarge)
1 node of instance type i3.4xlarge.

Instance ID	Public DNS (IPv4)
i-05e1d742bd7ece3e1	ec2-34-207-108-121.compute-1.amazonaws.com
Instance state	IPv4 Public IP

running	34.207.108.121
Instance type	IPv6 IPs
i3.4xlarge	-
Elastic IPs	
	Private DNS
1.	ip-172-31-39-224.ec2.internal
Availability zone	Private IPs
us-east-1c	172.31.39.224

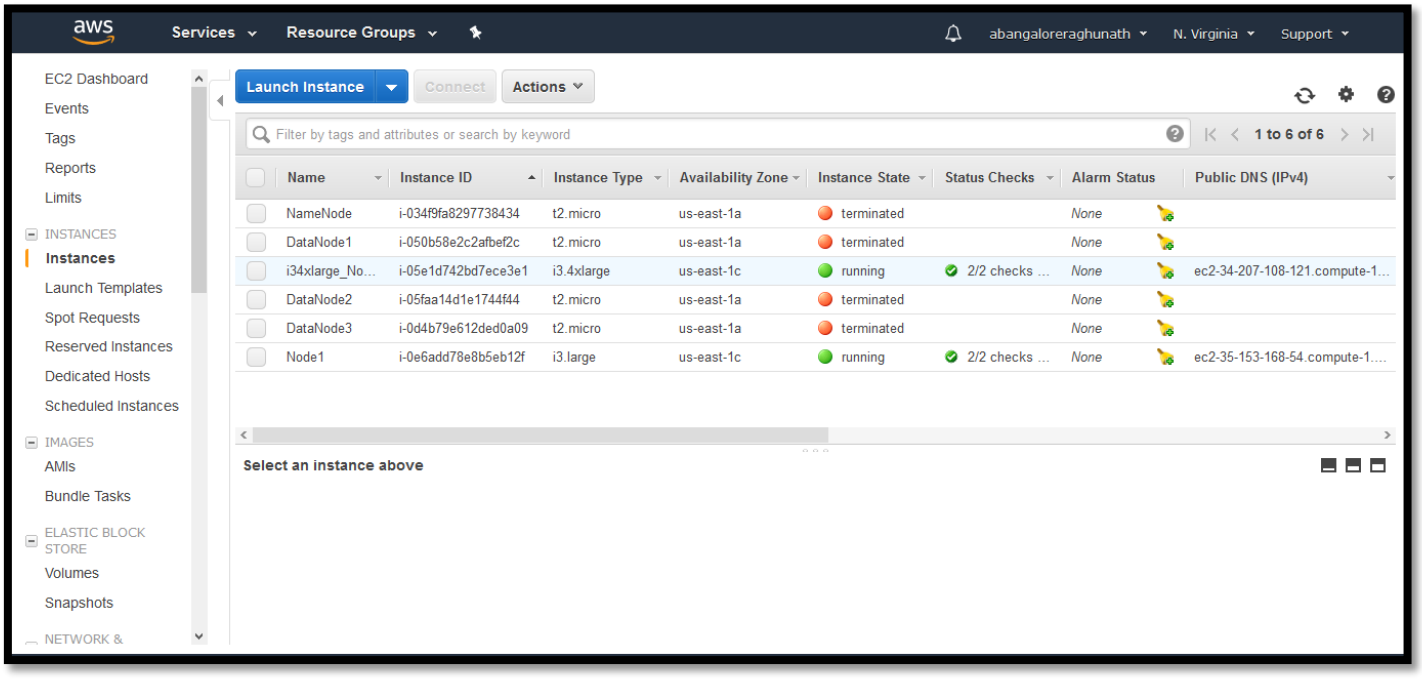

Shared memory

1TB

I faced a broken pipe exception twice, but finally it ran in the last run and completed the sort in 8 hours and 30 mins. 30,600 second.

Hadoop

1TB

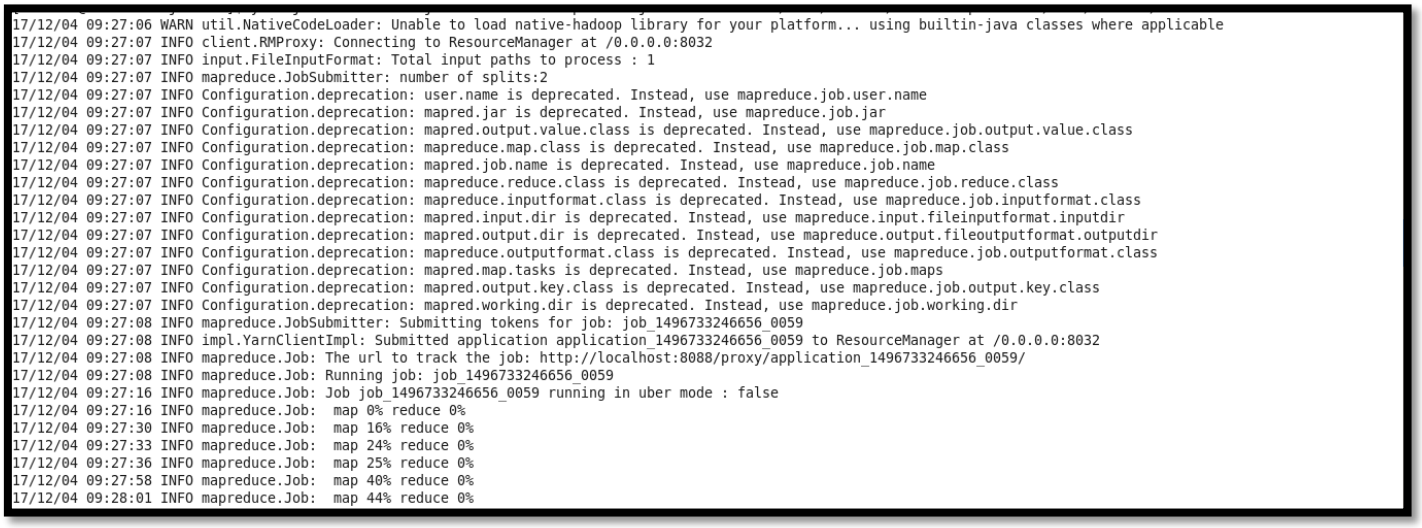

```
                Map output records=28006216
                Map output bytes=290097360
                Map output materialized bytes=346109804
                Input split bytes=234
                Combine input records=0
                Combine output records=0
                Reduce input groups=355
                Reduce shuffle bytes=346109804
                Reduce input records=28006216
                Reduce output records=355
                Spilled Records=84018648
                Shuffled Maps =2
                Failed Shuffles=0
                Merged Map outputs=2
                GC time elapsed (ms)=4240
                CPU time spent (ms)=45075030
                Physical memory (bytes) snapshot=412946432
                Virtual memory (bytes) snapshot=1082220544
                Total committed heap usage (bytes)=305799168
        Shuffle Errors
                BAD_ID=0
                CONNECTION=0
                IO_ERROR=0
                WRONG_LENGTH=0
                WRONG_MAP=0
                WRONG_REDUCE=0
        File Input Format Counters
                Bytes Read=1252156745864
        File Output Format Counters
                Bytes Written=23458692
SUCCESS
```

12:31 hours one node of i34x large

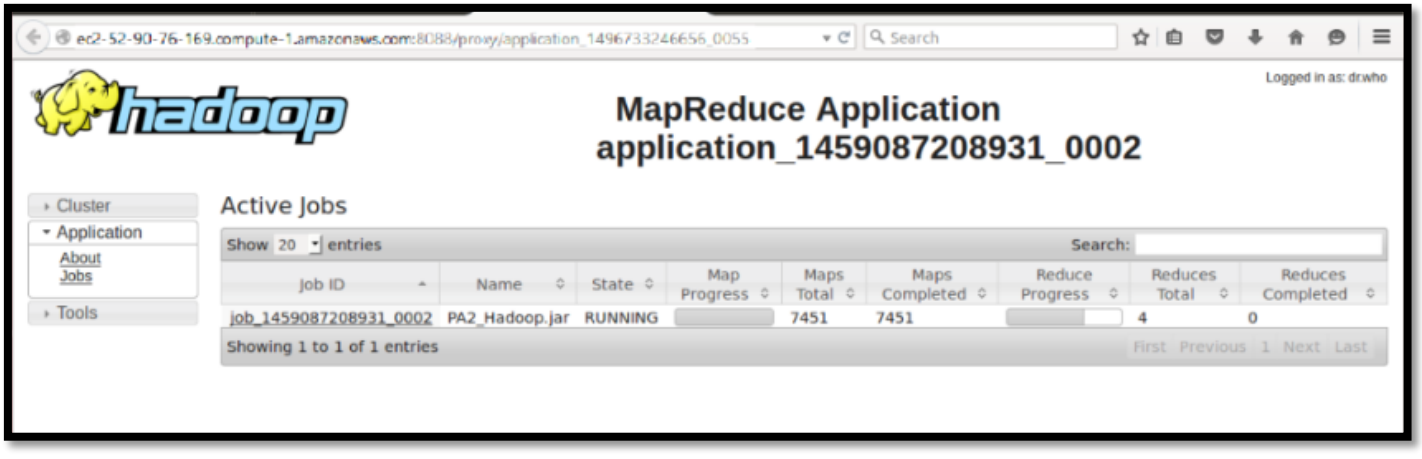

Spark

1TB i3.4xlarge instance virtual cluster Pseudo Mode.

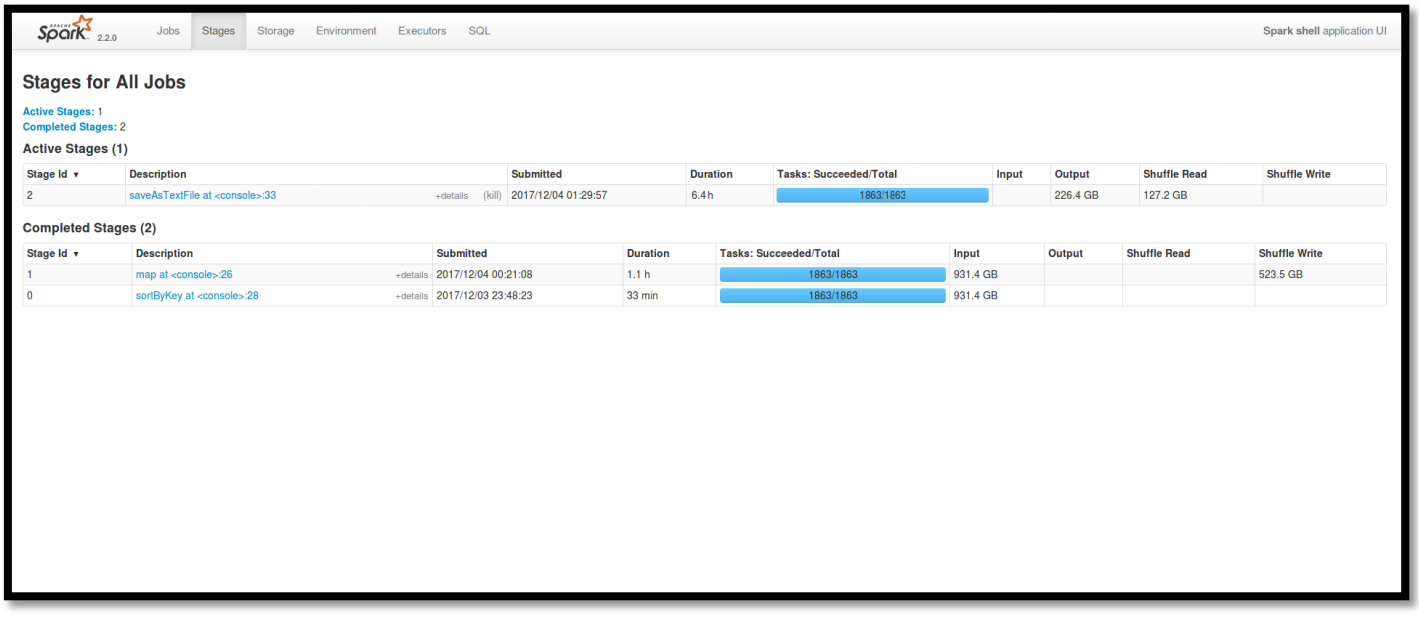

As we can see above the Spark instances could not complete initially, but they did eventually complete in **7.5** hours

Comparison of sorts on various levels of hardware quality and quantity

Experiment (instance/dataset)	Shared Memory TeraSort	Hadoop TeraSort	Spark TeraSort
Compute Time (sec) [1xi3.large 128GB]	13380	17688	11040
Data Read (GB) [1xi3.large 128GB]	134.08	131.256	128
Data Write (GB) [1xi3.large 128GB]	133.46	128.95	128
I/O Throughput (MB/sec) [1xi3.large 128GB]	19.98	14.71	23.188
Compute Time (sec) [1xi3.4xlarge 1TB]	30605	45060	27000
Data Read (GB) [1xi3.4xlarge 1TB]	1TB	1TB	1TB
Data Write (GB) [1xi3.4xlarge 1TB]	1TB	1TB	1TB
I/O Throughput (MB/sec) [1xi3.4xlarge 1TB]	65.39	44.38	74.07